

Byzantine Chain Replication

Pratik Sushil Zambani

pzambani@cs.stonybrook.edu, 111500248

Keshav Gupta

kegupta@cs.stonybrook.edu, 111464733

on receiving ..., on ... are handlers for specific events.
def func() are local functions.

Client

```
pending_requests = [] # tracks requests whose responses have not been
                        # received
replicas_public_keys = [] # to authenticate msgs from replicas
timer = Timer() # contains timer related functions like start, stop
```

on send request to head

```
request = client gets request from application
# sends a request to olympus for current configuration
config = olympus.get_config()
send ("request", request, request_id, retransmitted=0) to config.head
pending_requests.append(request_id)
# timer for the above request
timer.start()
```

on timer expire

```
# resend the request
config = olympus.get_config()
for each replica in config:
    send ("request", request, request_id, retransmitted=1) to replica
# timer for the above request
timer.start()
```

on receiving error statement:

```
# case when replica is immutable, fetch new config
config = olympus.get_config()
send ("request", request, request_id, retransmitted=1) to config.head
# timer for the above request
timer.start()
```

on receiving (result_struct)

```
if validate(result_struct.result_proof):
    timer.stop()
    pending_requests.remove(result_struct.result_proof.request_id)
    return result_struct.result to application
else:
    # send reconfiguration request to olympus with result_proof as proof of
    # misbehaviour
    send("reconfiguration_request", client, result_proof) to olympus
```

```
# runs periodically to check if configuration changed
on configuration check triggered (done periodically)
if config is not (updated_config = olympus.get_config()):
    config = updated_config
    for each request in pending_requests:
        execute sending request to head for request
```

```
on receiving("replicas_public_keys", replicas_public_keys) from Olympus
replicas_public_keys = replicas_public_keys
```

```
def validate(result, result_proof):
    # crypto_hash(arg) is a cryptographic hash function
    H = crypto_hash(result)
    count=0
    for replica_result_proof in result_proof:
        # check_signature validates signature
        if check_signature(replica_result_proof) and
            replica_result_proof.hash == H:
            count +=1
    if count < t+1
        return False
    return True
```

Replica

```
state_object = None
# mode is initlally PENDING
mode = PENDING          # in { ACTIVE, PENDING, IMMUTABLE }
history = []             # sequence of order proofs
result_cache = []
# cache to store result and result proof; limited size; uses FIFO strategy
timer = Timer(callback=send_reconfiguration_request)
crypto_keys = ()         # pair of public and private keys for self
replicas_public_keys = [] # public keys for all replicas
position = -1 # to distinguish between head, tail and internal replicas
last_checkpoint_proof = None # last checkpoint proof
slot = 0                 # points to current slot in history
```

```
on_receiving ("inithist", hist, running_state) from Olympus
history = hist
state_object = running_state
mode = ACTIVE
```

```
on_receiving ("key_pair", public_key, private_key) from Olympus
crypto_keys = (public_key, private_key)
```

```
on_receiving("replicas_public_keys", replicas_public_keys) from Olympus
replicas_public_keys = replicas_public_keys
```

```
# only called for head replica
on_receiving ("request", request, request_id, retransmitted=0)
slot = slot+1
# sign is a function to sign using the private key
if signed order_statements for all lower numbered slots present and
    mode == ACTIVE and history[slot] is None:
    order_statement = sign(<order, slot, request>, crypto_keys[1])
else:
    return error

result = state_object.evaluate_request(request)
# C is the configuration which the replica obtains from olympus
order_proof = (slot,request_id,request,replica,C,[order_statement])
history.append(order_proof)
# crypto_hash is a cryptographic hash function
result_proof = [sign(<"result", request_id, request, crypto_hash(result)>,
                    crypto_keys[1])]
shuttle = (order_proof , result_proof)
send (shuttle) to next replica
```

```

on receiving (shuttle) from predecessor position-1
# check for preconditions
if not is_order_proof_valid(shuttle.order_proof):
    # misbehaviour - different operations at a particular slot
    send("reconfiguration_request", replica) to olympus

# sign is a function to sign using the private key to send message to
# successor
if signed order_statements for all lower numbered slots present:
    order_statement = sign(<"order", slot, request>, crypto_keys[1])
else:
    send("reconfiguration_request", replica) to olympus
# We trigger a reconfiguration_request because if we ignore the request,
# client will retransmit, head has already ordered the operation and hence it #
will start a timer and wait, once timer expires reconfiguration will be
# triggered. So instead of delaying it, we trigger it right away.

request_id = shuttle.order_proof.request_id
result = state_object.evaluate_request(request)
shuttle.order_proof.order_stmt_list.append(order_statement)
shuttle.result_proof.append(sign(<"result", request_id, operation,
    crypto_hash(result)>), , crypto_keys[1])
history.append(shuttle.order_proof)

if position != 2t: # not tail
    send (shuttle) to next replica
else:
    send (result, shuttle.result_proof) to client
    result_cache[request_id] =
        (result, shuttle.result_proof)
    # send result shuttle back in the chain
    send (shuttle, result) to previous replica

```

```

on receiving (shuttle, result) from successor position+1
request_id = shuttle.result_proof.request_id
result_cache[request_id] = (result, shuttle.result_proof)
send (shuttle) to previous replica
# if timer is running for the request_id stop the timer and send the request
# to client
if timer(request_id):
    timer.stop()
    send (result_cache[request_id]) to client

```

```

on receiving ("request", request, request_id, retransmitted=1)
if is_correct(replica) and result_cache.contains(request_id):
    send (result_cache[request_id]) to client
else if mode == IMMUTABLE:
    return error
else if position != 0: # not head
    send ("request", request, request_id, retransmitted=1) to head
    timer.start(callback=send_reconfiguration_request)
else :
    # Head has seen the request and waiting for result proof from replicas
    if is_request_pending(request_id):
        timer.start(callback=send_reconfiguration_request)
    # treat it as a new request
    else:
        timer.start(callback=send_reconfiguration_request)
        execute on receiving ("request", request, request_id,
retransmitted=0)

```

```

on receiving ("wedge_request") from Olympus
mode = IMMUTABLE
# send history along with last checkpoint proof to olympus
send("wedged_statement", sign(<history, last_checkpoint_proof>,
                             crypto_keys[1]) to olympus

```

```

on receiving ("catch_up", missing_operations) from Olympus
for each operation in missing_operations:
    state_object.evaluate_request(operation)
send("caught_up", crypto_hash(state_object)) to olympus

```

```

on receiving ("get_running_state") from Olympus
send ("running_state", state_object) to olympus

```

```

# Checkpointing executed at head
on checkpointing triggered (done periodically after N new slots in history)
# state_object is the running state
checkpoint_statement = sign(<"checkpoint", crypto_hash(state_object)>,
                           crypto_keys[1])
checkpoint_proof_shuttle = [checkpoint_statement]
send (checkpoint_proof_shuttle, num_slots) to next replica

```

```

on receiving (checkpoint_proof_shuttle, num_slots) from predecessor
checkpoint_statement = sign(<"checkpoint", crypto_hash(state_object)>,
    crypto_keys[1])
checkpoint_proof_shuttle.append(checkpoint_statement)

# at tail, checkpoint proof is completed and sent back in the chain
if position == 2t:
    completed_checkpoint_proof = checkpoint_proof_shuttle
    if are_signatures_valid(completed_checkpoint_proof):
        # truncates oldest num_slots from history
        history.truncate(num_slots)
        # save last_checkpoint_proof
        last_checkpoint_proof = completed_checkpoint_proof

    send (completed_checkpoint_proof, num_slots) to predecessor replica

# if not tail, send checkpoint shuttle forward in the chain
else:
    send (checkpoint_proof_shuttle, num_slots) to next replica



---


on receiving (completed_checkpoint_proof, num_slots) from successor
# everybody in chain has received checkpoint msg, start truncating
if are_signatures_valid(completed_checkpoint_proof):
    # truncates oldest num_slots from history
    history.truncate(num_slots)
    # save last_checkpoint_proof
    last_checkpoint_proof = completed_checkpoint_proof

if pos != 0: # not head
    send (checkpoint_proof_shuttle, num_slots) to predecessor replica



---


# callback for timer expire
def send_reconfiguration_request():
    send ("reconfiguration_request", replica) to olympus



---


def are_signatures_valid(proof):
    for each rho, order_statement in proof:
        key = replicas_public_keys[rho]
        if not check_signature(order_statement, key)
            return False

    return True



---



```

```

# Checks for valid signatures and operation in order proof
def is_order_proof_valid(order_proof):
    operation = order_proof.operation
    for each rho, order_statement in order_proof:
        key = replicas_public_keys[rho]
        if not check_signature(order_statement, key) or
            order_statement.operation is not operation:
            return False
    return True



---


def is_correct(replica):
    return replica.mode in (PENDING, ACTIVE, IMMUTABLE) and
        is_history_valid(replica.history)



---


def is_history_valid(hist):
    for each slot in hist:
        if not all(hist.order_proofs.operation) == hist.operation:
            return False
    return True

# checks if request is received and result is not available yet to be sent to
# client because result shuttle hasn't been received



---


def is_request_pending(request_id):
    if request_id in any(history.order_proof) and
        request_id not in result_cache:
        return True
    return False



---



```


Olympus

```
# C is current configuration, C.head used in pseudocode code is C[0],
# C.tail is C[-1]...
C = [replica identifiers ...]
clients = [client identifiers ...] # list of clients
keys = [] #list of tuples that stores public and private keys for each replica
```

```
on receiving ("reconfiguration_request", sender, proof_of_misbehaviour=None)
```

```
# if sender is not in current configuration and the proof of misbehaviour is
# invalid, ignore
if sender not in C and (proof_of_misbehaviour == None or not
    check_proof_of_misbehaviour(proof_of_misbehaviour)):
    return # ignore request
```

```
for each replica in C:
    send("wedge_request") to replica
```

```
# receiving wedged statements
await atleast t+1 consistent wedged statements from replicas and store in
wedged_statements
```

```
valid_quorum = False
while not valid_quorum:
```

```
    quorum_candidate = select any t+1 wedged_statements
```

```
    # Q is quorum
    # condition for consistency: for every pair of wedged statements,
    # w1, w2 for each slot operation should be same in w1.history and
    # w2.history. Signature of the messages should be valid. Also
    # wedged_statement.checkpoint should be valid for quorum of replicas.
    if is_history_consistent(quorum_candidate):
        Q = quorum_candidate
    else:
        continue
```

```
    # longest history - most number of order proofs
    LH = select longest history order proof from Q
    for each replica in Q:
        send("catch_up", LH-replica.wedged_statement.history) to replica
```

```
    await |Q| "caught_up" messages from replicas in caught_up_messages
    valid_quorum = True
```

```

    # check if hash in caught_up_messages are same
    hash_val = any(caught_up_msgs.hash)
    for each msg in caught_up_msgs:
        if not msg.hash == hash_val:
            valid_quorum = False
            break

for each replica in Q:
    send("get_running state") to replica
    await "running_state" message from replica into running_state
    if(crypto_hash(running_state) == hash_val):
        break

# start 2t+1 new replicas in new configuration C'
for i in xrange(2t+1):
    # replica_i is identifier for replica
    replica = Replica(replica_i)
    replica.start()
    C'.append(replica)

for each replica in C':
    # the new history on the replica will be empty
    new_hist = []
    send ("inithist", new_hist, running_state) to replica
    # generate_keys generates public key cryptography keys
    (public_key, private_key) = generate_keys(replica)
    send ("key_pair", public_key, private_key) to replica
    keys[replica.replica_id] = (public_key, private_key)

# send all public keys of others to all replicas and clients
replicas_public_keys = get all public_keys from keys
for each replica in C':
    send ("replicas_public_keys", replicas_public_keys) to replica
for each client in clients:
    send ("replicas_public_keys", replicas_public_keys) to client
C = C'

```

```

def check_proof_of_misbehaviour(result_proof):
    H = result_proof[0].hash
    # if signature is not valid or hash of any result proof mismatach it
    # indicates a proof of misbehaviour
    for replica_result_proof in result_proof:
        if not (check_signature(replica_result_proof) and
                replica_result_proof.hash == H):
            return False

    return True

```

```

def is_history_consistent(wedged_statements):
    # check consistency between every pair of wedged statements
    for every w1 in wedged_statements:
        for every w2 in wedged_statements:
            if w1 == w2:
                continue
            for slot in range(0, min(w1.history.length,
                                     w2.history.length)):
                if w1.history[slot].request !=
                   w2.history[slot].request:
                    return False

    # check validity of checkpoint proof
    hash_val = wedged_statements[0].last_checkpoint_proof.hash
    for each last_checkpoint_proof in wedged_statements:
        if last_checkpoint_proof.hash != hash_val
            return False

    return True

```

```

def get_config():
    return C

```
